

COVID19_Analysis

Qizheng Wang

4/10/2020

```
library(devtools)
```

```
## Loading required package: usethis
```

```
library(mgcv)
```

```
## Loading required package: nlme
```

```
## This is mgcv 1.8-31. For overview type 'help("mgcv-package")'.
```

```
library(gamm4)
```

```
## Loading required package: Matrix
```

```
## Loading required package: lme4
```

```
## Warning: package 'lme4' was built under R version 3.6.3
```

```
##
```

```
## Attaching package: 'lme4'
```

```
## The following object is masked from 'package:nlme':
```

```
##
```

```
##      lmList
```

```
## This is gamm4 0.2-6
```

```
library(tidyverse)
```

```
## -- Attaching packages -----
```

```
## v ggplot2 3.2.1    v purrr  0.3.3
```

```
## v tibble  2.1.3    v dplyr  0.8.4
```

```
## v tidyr   1.0.0    v stringr 1.4.0
```

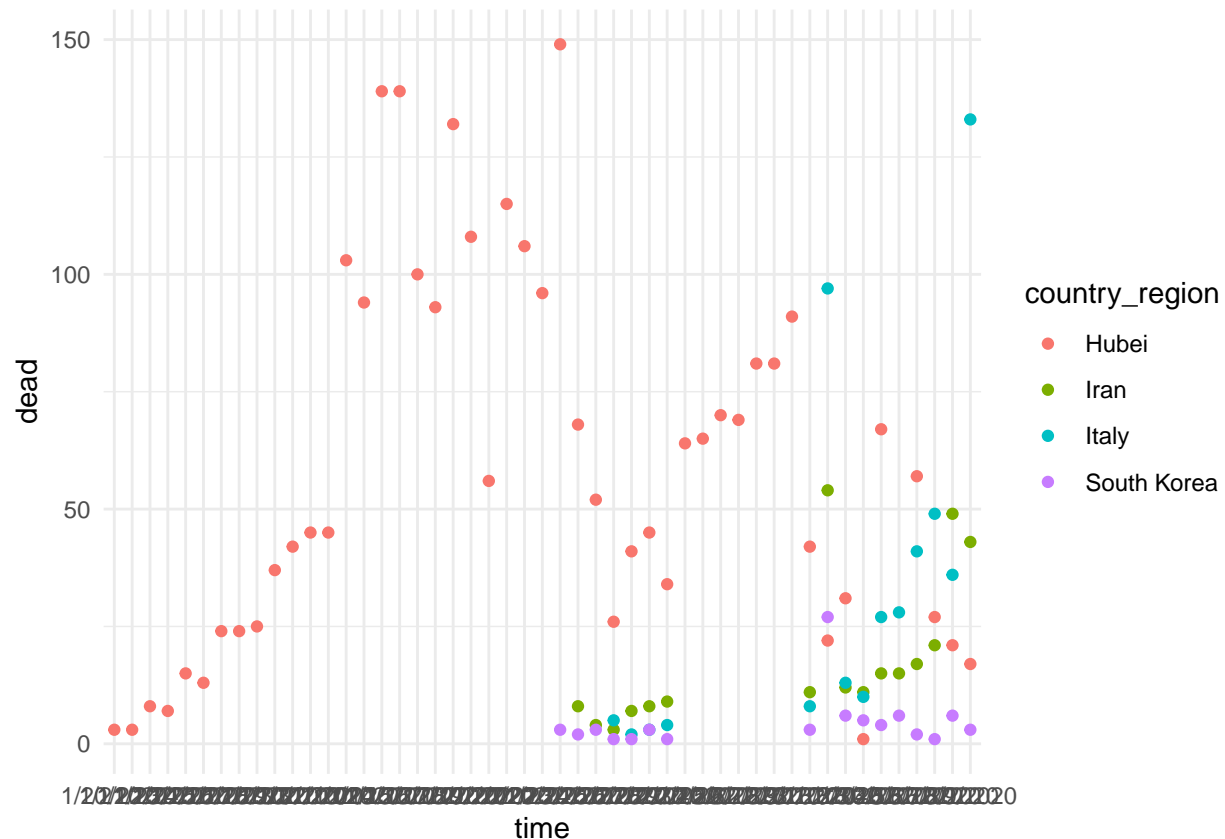
```
## v readr   1.3.1    v forcats 0.4.0
```

```
## -- Conflicts ----- tidyverse
## x dplyr::collapse() masks nlme::collapse()
## x tidyr::expand() masks Matrix::expand()
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
## x tidyr::pack() masks Matrix::pack()
## x tidyr::unpack() masks Matrix::unpack()
```

```
setwd("c:/Users/roder/OneDrive/Desktop/COVID")
covid_data <- read.csv("covid_data.csv")
```

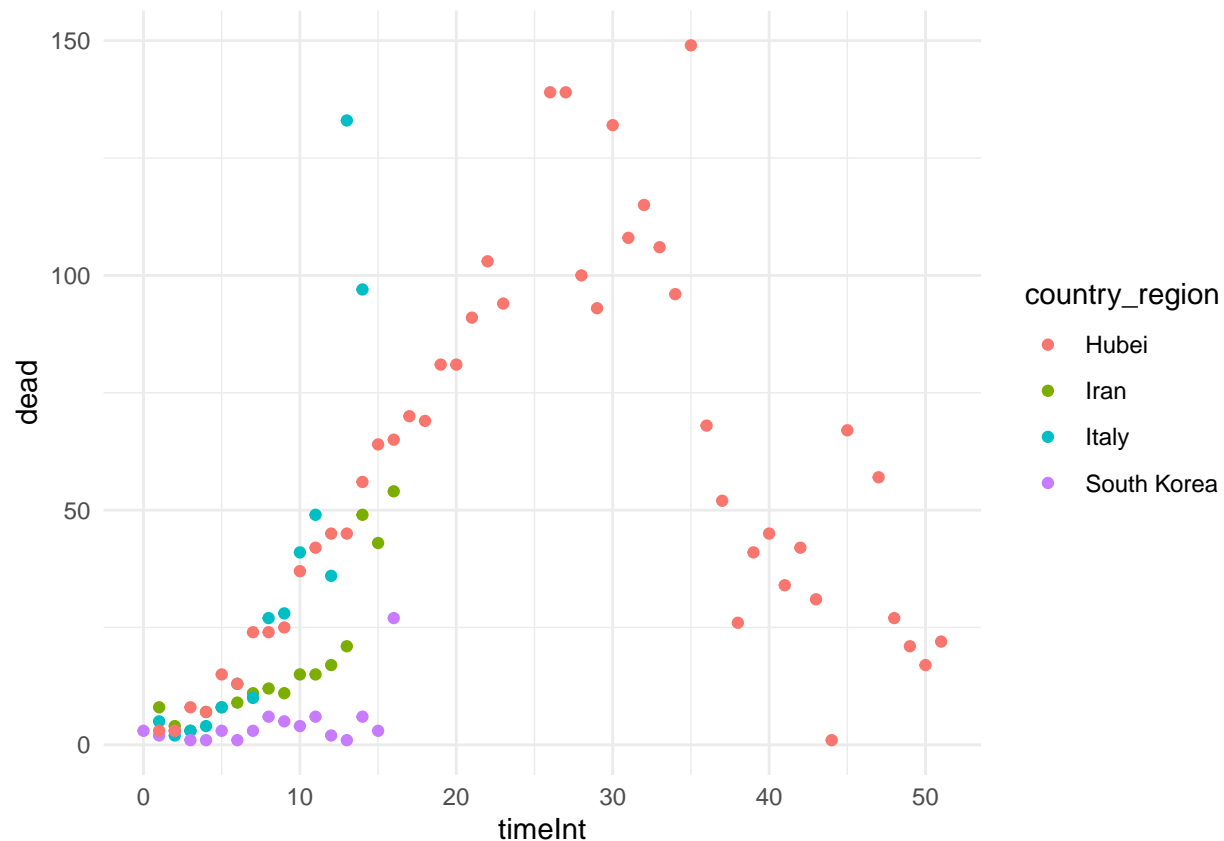
```
# Plot over time
```

```
covid_data %>%
  filter(country_region %in% c('Hubei', 'Italy', 'Iran', 'South Korea', 'USA')) %>%
  na.omit() %>%
  ggplot(aes(time, dead, color=country_region)) +
  geom_point() +
  theme_minimal()
```



```
# Plot from initial death in region
```

```
covid_data %>%
  filter(country_region %in% c('Hubei', 'Italy', 'Iran', 'South Korea', 'USA')) %>%
  na.omit() %>%
  ggplot(aes(timeInt, dead, color=country_region)) +
  geom_point() +
  theme_minimal()
```



```
## Setting up GAM Model
## Because timeInt indicates date since the first death, so we should constrain the line to pass
## through the origin when timeInt = 0, indicating there are no death before any deaths occurred

resGam= mgcv::gam(
  dead ~ s(timeInt, pc=0) + country_region,
  data=covid_data,
  family=poisson(link='log'))

summary(resGam)
```

```
##
## Family: poisson
## Link function: log
##
## Formula:
## dead ~ s(timeInt, pc = 0) + country_region
##
## Parametric coefficients:
##
```

	Estimate	Std. Error	z value	Pr(> z)
## (Intercept)	-0.160352	0.583136	-0.275	0.783329
## country_regionAustralia	0.078106	1.155196	0.068	0.946094
## country_regionBeijing	-1.940556	0.739512	-2.624	0.008688 **
## country_regionChongqing	-0.535153	0.819679	-0.653	0.513833
## country_regionFrance	1.127419	0.610845	1.846	0.064940 .
## country_regionGuangdong	-1.608135	0.771882	-2.083	0.037215 *

```

## country_regionHainan      -2.168937    0.824279   -2.631 0.008506 **
## country_regionHebei       -0.763389    0.823787   -0.927 0.354092
## country_regionHeilongjiang -1.118993    0.666038   -1.680 0.092943 .
## country_regionHenan       -1.208796    0.631050   -1.916 0.055425 .
## country_regionHubei        1.815819    0.589066    3.083 0.002052 **
## country_regionHunan        0.078106    1.155196    0.068 0.946094
## country_regionIran         1.321243    0.590201    2.239 0.025180 *
## country_regionIraq         0.171690    0.764797    0.224 0.822375
## country_regionItaly        2.117238    0.588802    3.596 0.000323 ***
## country_regionJapan       -1.361864    0.654921   -2.079 0.037578 *
## country_regionShandong     0.215099    0.817422    0.263 0.792440
## country_regionSouth Korea  -0.005497    0.597876   -0.009 0.992664
## country_regionSpain        2.033865    0.605583    3.359 0.000784 ***
## country_regionUnited Kingdom 1.258965    0.820598    1.534 0.124979
## country_regionUnited States 0.827315    0.621365    1.331 0.183042
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
## s(timeInt) 8.758  8.982   1309  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.894   Deviance explained = 93.5%
## UBRE = 2.0019   Scale est. = 1           n = 170

```

```
coef(resGam)
```

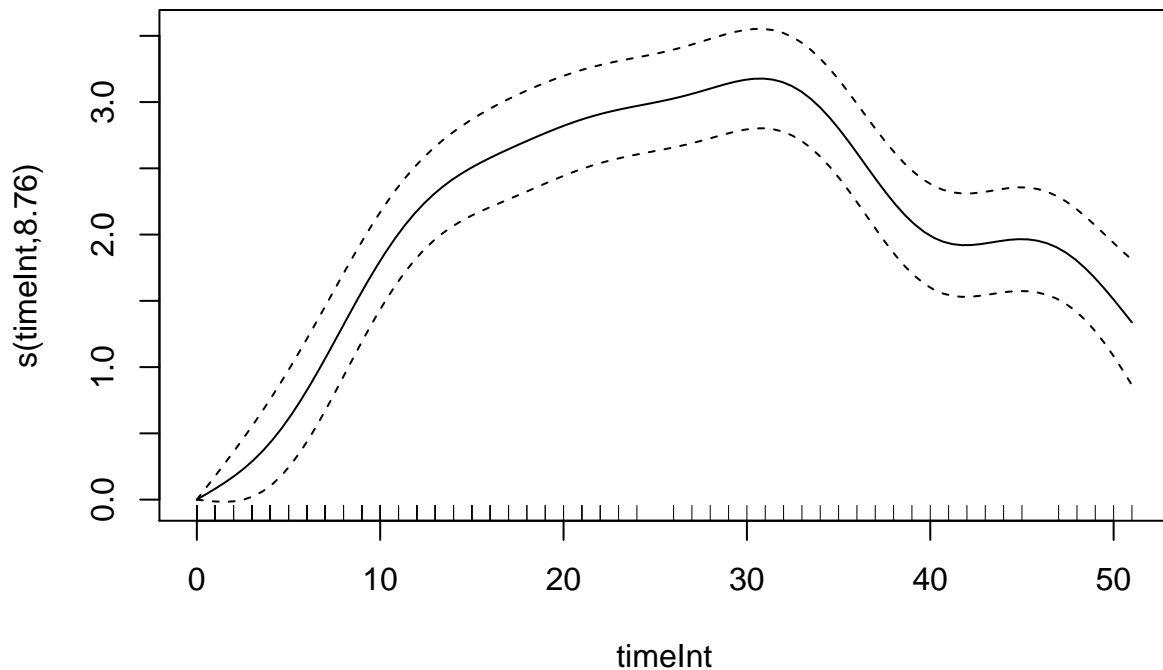
```

##              (Intercept)      country_regionAustralia
##              -0.160352456           0.078105515
## country_regionBeijing      country_regionChongqing
##              -1.940556292           -0.535153159
## country_regionFrance      country_regionGuangdong
##              1.127419488           -1.608135374
## country_regionHainan      country_regionHebei
##              -2.168937066           -0.763389041
## country_regionHeilongjiang  country_regionHenan
##              -1.118993096           -1.208796089
## country_regionHubei      country_regionHunan
##              1.815818734           0.078105515
## country_regionIran      country_regionIraq
##              1.321243223           0.171690309
## country_regionItaly      country_regionJapan
##              2.117237701           -1.361864231
## country_regionShandong    country_regionSouth Korea
##              0.215099168           -0.005496802
## country_regionSpain      country_regionUnited Kingdom
##              2.033864959           1.258964745
## country_regionUnited States s(timeInt).1
##              0.827314895           0.436070190
##              s(timeInt).2           s(timeInt).3
##              0.162668721           0.695995274
##              s(timeInt).4           s(timeInt).5

```

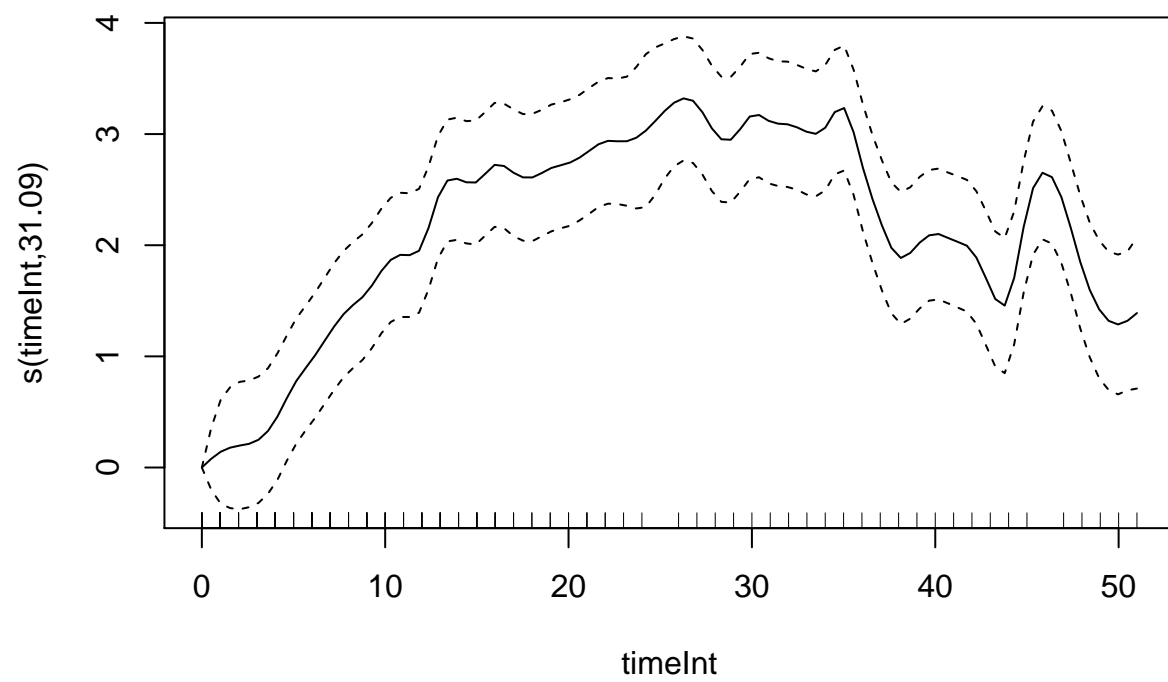
```
##          -0.257405570          0.133254518
##          s(timeInt).6          s(timeInt).7
##          1.140898783          -0.022139449
##          s(timeInt).8          s(timeInt).9
##          4.992873514          -1.020359041
```

```
plot(resGam)
```

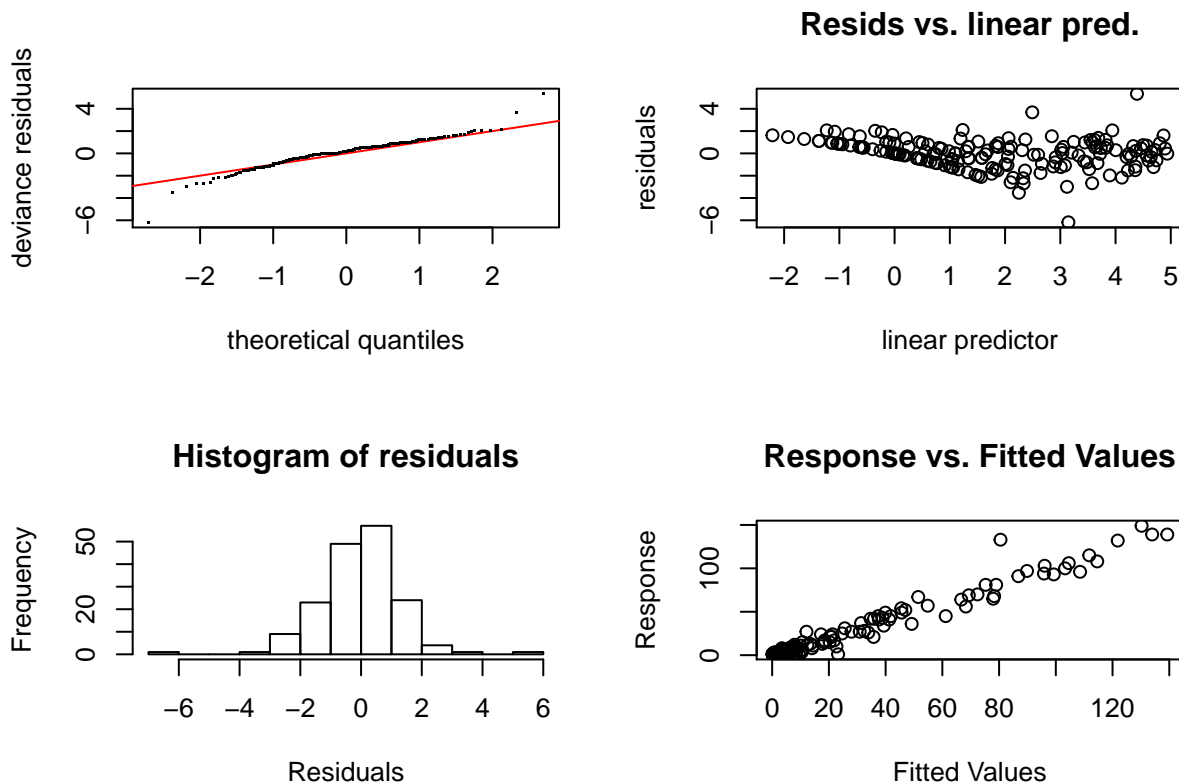


Analysis: the estimated degrees of freedom for the smooth of timeInt is 8.758, Since edf is much higher than 1, the relationship is not linear. And we can't interpret the coefficients for the smooth of timeInt because they are coefficients for the different splines that make up our curve, but don't have a scientific interpretation, thus we cannot interpret the coefficients for country_region as usual.

```
resGam3= mgcv::gam(
dead ~ s(timeInt, k=50, pc=0) + country_region, data=covid_data,
family=poisson(link='log'), method='ML')
plot(resGam3)
```

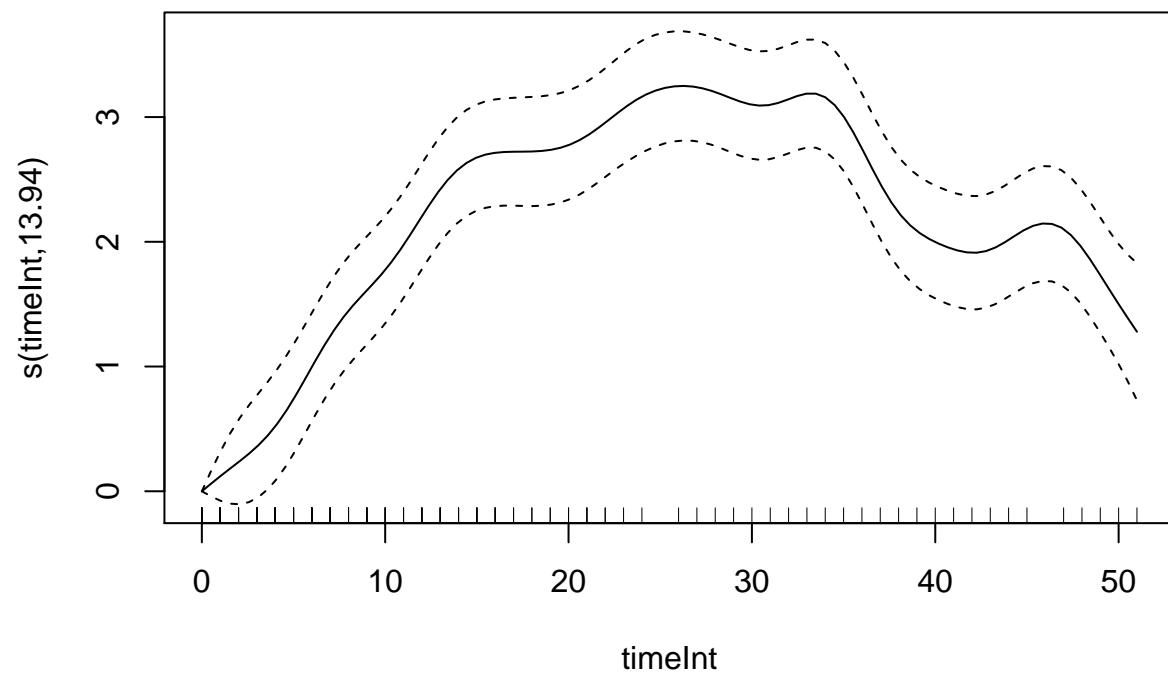


```
gam.check(resGam3)
```

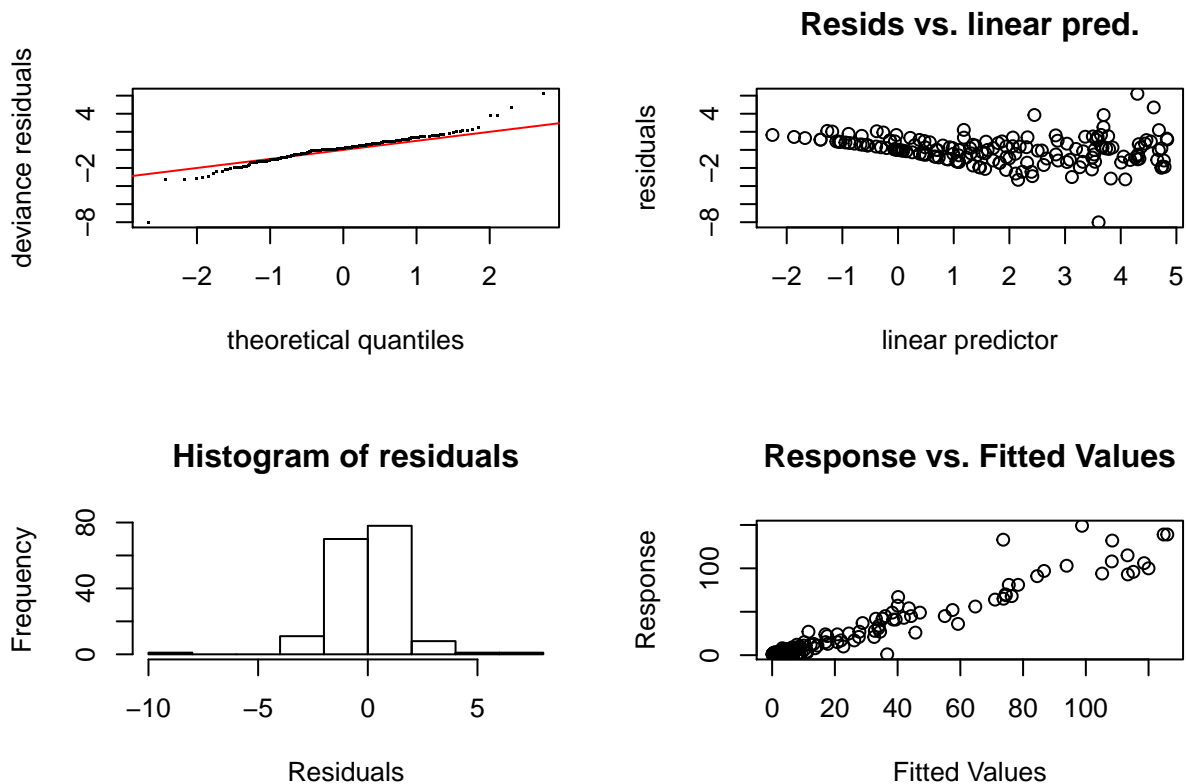


```
##
## Method: ML    Optimizer: outer newton
## full convergence after 6 iterations.
## Gradient range [-1.704072e-05,-1.704072e-05]
## (score 540.3471 & scale 1).
## Hessian positive definite, eigenvalue range [4.080029,4.080029].
## Model rank = 70 / 70
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##           k'   edf k-index p-value
## s(timeInt) 49.0 31.1   1.25      1
```

```
resGam4 = mgcv::gam(
  dead ~ s(timeInt, k=20, pc=0) + country_region, data=covid_data,
  family=poisson(link='log'), method='ML')
plot(resGam4)
```



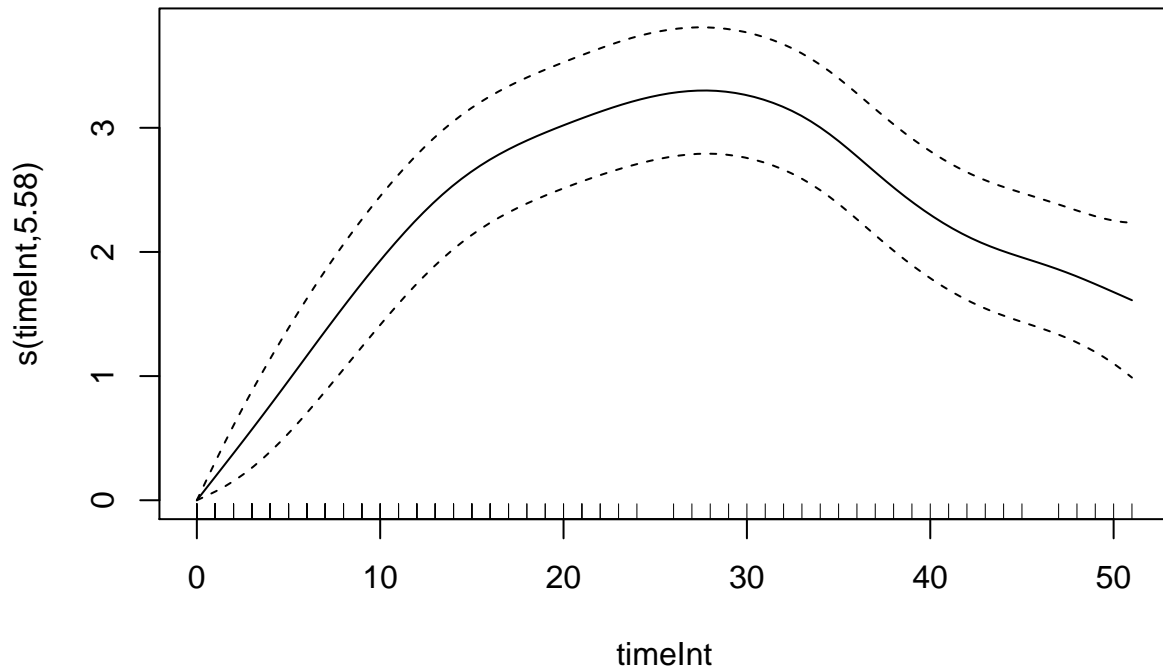
```
gam.check(resGam4)
```

```
##
## Method: ML   Optimizer: outer newton
## full convergence after 6 iterations.
## Gradient range [3.691928e-06,3.691928e-06]
## (score 554.3095 & scale 1).
## Hessian positive definite, eigenvalue range [3.724135,3.724135].
## Model rank = 40 / 40
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##           k'   edf k-index p-value
## s(timeInt) 19.0 13.9   1.15   0.95
```

In this part, choose different k , $k = 50$ and $k = 20$. And run `gam.check()` for both of them. After seeing the result, I would choose $k = 20$ for capturing patterns in the data without overfitting the data. However data seems highly correlated. A random effect for the country should be fitted.

```
covid_data$timeIntInd = covid_data$timeInt
resGammInd = gamm4::gamm4(
  dead ~ country_region +
  s(timeInt, k=20, pc=0),
  random = ~ (1|timeIntInd),
  data=covid_data, family=poisson(link='log'))
plot(resGammInd$gam)
```



```
summary(resGammInd$mer)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: poisson ( log )
##
##      AIC      BIC   logLik deviance df.resid
##  1082.2   1157.4   -517.1  1034.2     146
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.2542 -0.5002  0.0522  0.8694  5.2818
##
## Random effects:
##   Groups      Name      Variance Std.Dev.
## timeIntInd (Intercept) 0.08203  0.2864
## Xr          s(timeInt) 5.19007  2.2782
## Number of obs: 170, groups:  timeIntInd, 50; Xr, 18
##
## Fixed effects:
##
##              Estimate Std. Error z value Pr(>|z|)
## X(Intercept)    -0.306480   0.605145  -0.506  0.612536
## Xcountry_regionAustralia    0.006297   1.163527   0.005  0.995682
## Xcountry_regionBeijing    -2.011547   0.741361  -2.713  0.006661 **
## Xcountry_regionChongqing   -0.656657   0.823396  -0.797  0.425162
```

```
## Xcountry_regionFrance      1.045405    0.612873    1.706 0.088055 .
## Xcountry_regionGuangdong   -1.641456    0.775393   -2.117 0.034265 *
## Xcountry_regionHainan      -2.299227    0.843758   -2.725 0.006430 **
## Xcountry_regionHebei       -0.882402    0.825837   -1.068 0.285298
## Xcountry_regionHeilongjiang -1.054884    0.668823   -1.577 0.114744
## Xcountry_regionHenan       -1.241604    0.633073   -1.961 0.049852 *
## Xcountry_regionHubei       1.772212    0.590969    2.999 0.002710 **
## Xcountry_regionHunan       0.006266    1.163493    0.005 0.995703
## Xcountry_regionIran        1.236439    0.592258    2.088 0.036828 *
## Xcountry_regionIraq        0.151040    0.768669    0.196 0.844223
## Xcountry_regionItaly       2.044959    0.590769    3.462 0.000537 ***
## Xcountry_regionJapan       -1.417716    0.657006   -2.158 0.030940 *
## Xcountry_regionShandong     0.083982    0.822740    0.102 0.918697
## Xcountry_regionSouth Korea -0.088619    0.599821   -0.148 0.882546
## Xcountry_regionSpain       2.018165    0.604940    3.336 0.000850 ***
## Xcountry_regionUnited Kingdom 1.338358    0.832920    1.607 0.108093
## Xcountry_regionUnited States 0.745296    0.623271    1.196 0.231782
## Xs(timeInt)Fx1            2.801119    0.765139    3.661 0.000251 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Correlation matrix not shown by default, as p = 22 > 12.
## Use print(x, correlation=TRUE) or
##     vcov(x)           if you need it
```

```
summary(resGammInd$gam)
```

```
##
## Family: poisson
## Link function: log
##
## Formula:
## dead ~ country_region + s(timeInt, k = 20, pc = 0)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.306480   0.608607  -0.504 0.614559
## country_regionAustralia    0.006297   1.169943   0.005 0.995705
## country_regionBeijing     -2.011547   0.744893  -2.700 0.006925 **
## country_regionChongqing   -0.656657   0.827599  -0.793 0.427517
## country_regionFrance      1.045405   0.616690   1.695 0.090040 .
## country_regionGuangdong   -1.641456   0.779151  -2.107 0.035141 *
## country_regionHainan      -2.299227   0.850216  -2.704 0.006845 **
## country_regionHebei       -0.882402   0.830000  -1.063 0.287721
## country_regionHeilongjiang -1.054884   0.672528  -1.569 0.116756
## country_regionHenan       -1.241604   0.636748  -1.950 0.051186 .
## country_regionHubei       1.772212   0.594628   2.980 0.002879 **
## country_regionHunan       0.006266   1.169956   0.005 0.995727
## country_regionIran        1.236439   0.595903   2.075 0.037996 *
## country_regionIraq        0.151040   0.773174   0.195 0.845119
## country_regionItaly       2.044959   0.594420   3.440 0.000581 ***
## country_regionJapan       -1.417716   0.660632  -2.146 0.031873 *
```

```
## country_regionShandong      0.083982    0.827671    0.101 0.919180
## country_regionSouth Korea   -0.088619    0.603459   -0.147 0.883249
## country_regionSpain         2.018165    0.608779    3.315 0.000916 ***
## country_regionUnited Kingdom 1.338358    0.839214    1.595 0.110762
## country_regionUnited States  0.745296    0.627131    1.188 0.234667
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
## s(timeInt)  5.58   5.58  289.7  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.884
## glmer.ML = 250.06  Scale est. = 1          n = 170
```

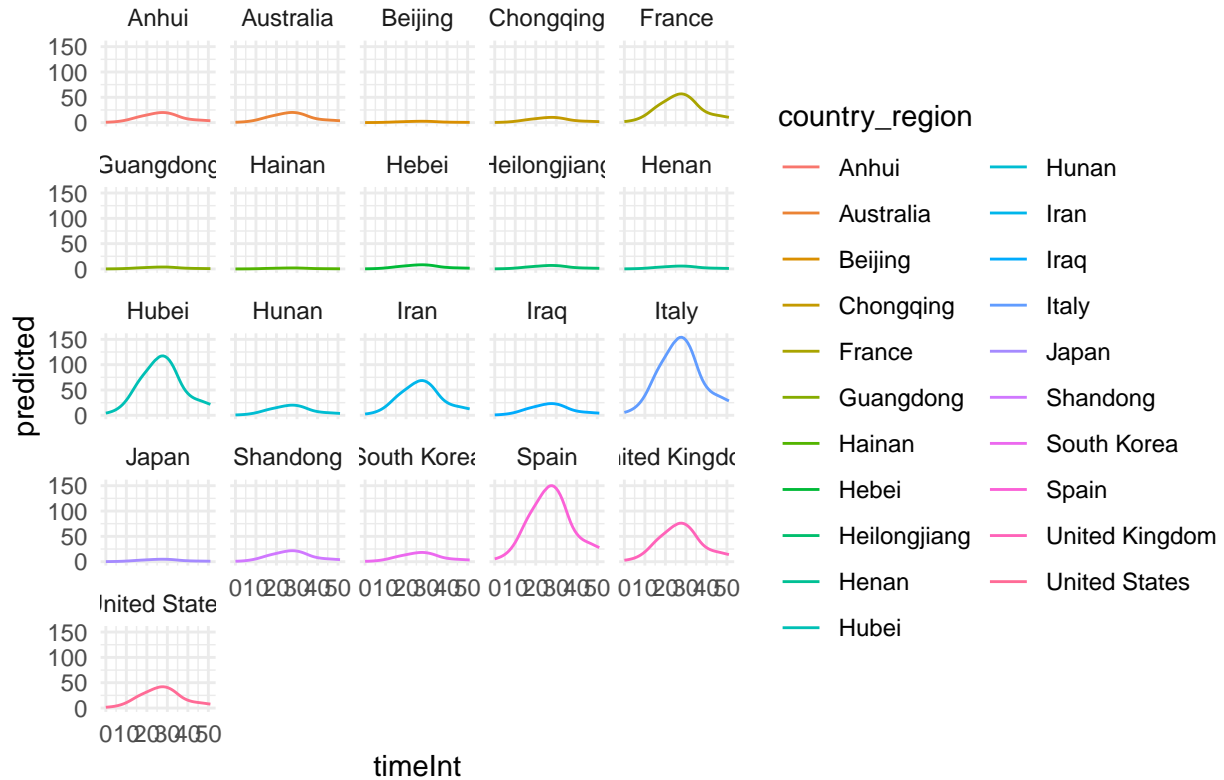
So use counry_region nessted within timIntInd to fit another same model.

This suggests a tren where there is a shaper increase in the deaths per day over the first 25 days to a month and the number decreases the following 30 days.

```
covid_data_2 <- expand_grid(covid_data$timeInt, covid_data$country_region) %>%
as_tibble() %>%
rename(timeInt = 1, country_region = 2) %>%
distinct()
covid_data_2$predicted <- predict(resGammInd$gam, newdata=covid_data_2, type="response")

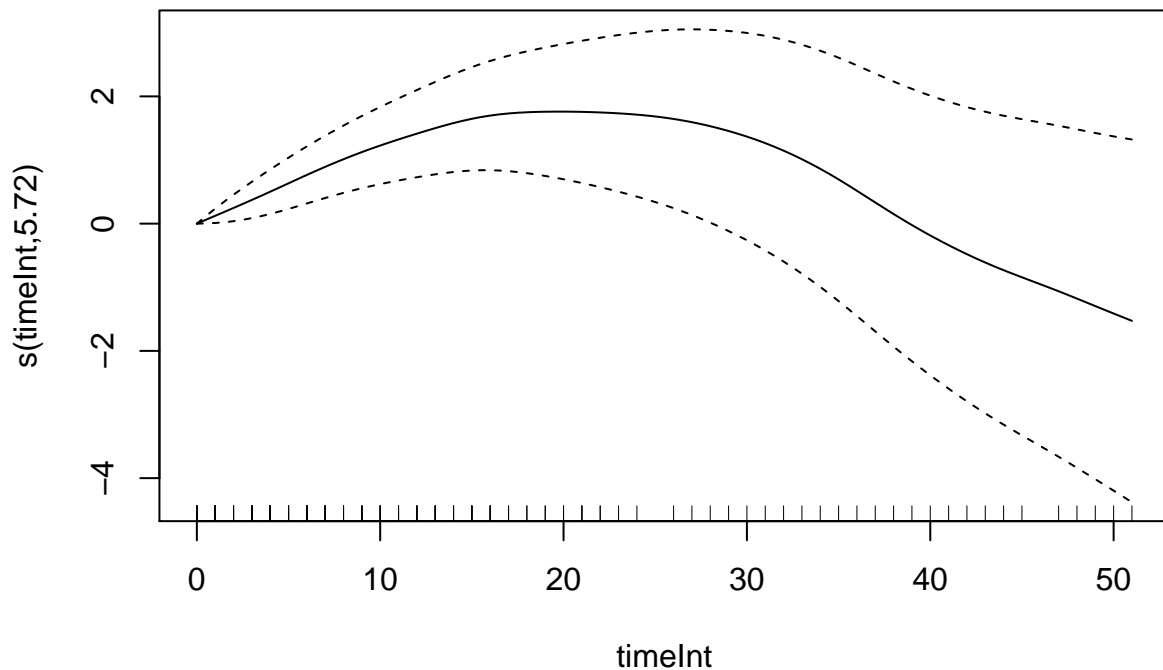
covid_data_2 %>%
ggplot(aes(timeInt, predicted, colour=country_region)) +
geom_line() +
theme_minimal() +
facet_wrap(~country_region) +
ggtitle("Predicted deaths over time (time = 0 is first death)")
```

Predicted deaths over time (time = 0 is first death)



```
covid_data$timeSlope = covid_data$timeInt/100
resGammSlope = gamm4::gamm4(
  dead ~ country_region + s(timeInt, k=30, pc=0),
  random = ~(0+timeSlope|country_region) +
  (1|timeIntInd:country_region),
  data=covid_data, family=poisson(link='log'))

plot(resGammSlope$gam)
```



```
summary(resGammSlope$mer)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
##   Approximation) [glmerMod]
##   Family: poisson ( log )
##
##      AIC      BIC   logLik deviance df.resid
##    991.2   1069.6  -470.6   941.2     145
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.2172 -0.3074 -0.0140  0.2211  2.0847
##
## Random effects:
##   Groups                Name             Variance Std.Dev.
## timeIntInd:country_region (Intercept)  0.08517  0.2918
## Xr                                s(timeInt)  3.57360  1.8904
## country_region              timeSlope    55.12944  7.4249
## Number of obs: 170, groups:
## timeIntInd:country_region, 170; Xr, 28; country_region, 21
##
## Fixed effects:
##
##              Estimate Std. Error z value Pr(>|z|)
## X(Intercept)   -0.24988    0.62034  -0.403  0.68709
## Xcountry_regionAustralia  0.09264    1.20949   0.077  0.93894
```

```
## Xcountry_regionBeijing      -0.63006    1.20219   -0.524    0.60021
## Xcountry_regionChongqing    -0.25543    0.92027   -0.278    0.78135
## Xcountry_regionFrance       0.95325    0.69546    1.371    0.17048
## Xcountry_regionGuangdong    -0.31269    0.95823   -0.326    0.74418
## Xcountry_regionHainan       -0.56894    1.18326   -0.481    0.63065
## Xcountry_regionHebei        -0.55685    0.98029   -0.568    0.57000
## Xcountry_regionHeilongjiang  0.14675    0.77880    0.188    0.85053
## Xcountry_regionHenan        0.43946    0.72421    0.607    0.54397
## Xcountry_regionHubei        1.80314    0.65407    2.757    0.00584 **
## Xcountry_regionHunan        0.09280    1.20941    0.077    0.93884
## Xcountry_regionIran         1.34649    0.66567    2.023    0.04310 *
## Xcountry_regionIraq         0.16359    0.83447    0.196    0.84458
## Xcountry_regionItaly        0.98696    0.68097    1.449    0.14724
## Xcountry_regionJapan        0.17602    0.82260    0.214    0.83056
## Xcountry_regionShandong     0.24189    0.89838    0.269    0.78774
## Xcountry_regionSouth Korea   0.46834    0.69393    0.675    0.49974
## Xcountry_regionSpain        1.97649    0.66869    2.956    0.00312 **
## Xcountry_regionUnited Kingdom 1.31487    0.89625    1.467    0.14235
## Xcountry_regionUnited States 0.93941    0.69632    1.349    0.17730
## Xs(timeInt)Fx1              1.57790    0.81609    1.933    0.05318 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Correlation matrix not shown by default, as p = 22 > 12.
## Use print(x, correlation=TRUE) or
##     vcov(x)           if you need it
```

```
names(lme4::ranef(resGammSlope$mer))
```

```
## [1] "timeIntInd:country_region" "Xr"
## [3] "country_region"
```

```
theRanef = lme4::ranef(resGammSlope$mer, condVar = TRUE)$country_region
(theRanefVec = sort(drop(t(theRanef))))
```

```
##           Japan           Henan    Heilongjiang      Guangdong      Hainan
##    -7.45640077    -7.39197752    -6.59756826    -4.01398844    -3.18040196
##           Beijing  United States      Chongqing           Anhui           Hebei
##    -2.65471706    -1.74396156    -1.45000907    -0.17970022    -0.15082900
##           Iraq  United Kingdom           Hunan      Australia      Shandong
##    -0.02958292     0.00000000     0.01712546     0.01720012     0.25261765
##    South Korea           Spain           France           Iran           Hubei
##     1.40347325     3.16651164     5.55065837     5.63517909     6.01194601
##           Italy
##    16.14488005
```

```
Dcountry = 'France'
toPredict = expand.grid(
  timeInt = 0:100,
  country_region = Dcountry)
toPredict$timeSlope = toPredict$timeIntInd =
```

```

toPredict$timeInt

thePred = predict(resGammSlope$gam,
newdata=toPredict, se.fit=TRUE)
matplot(toPredict$timeInt,
exp(do.call(cbind, thePred) %*% Pmisc::ciMat(0.75)),
type='l',
col=c('black','grey','grey'),
ylim = c(0, 25))
points(covid_data[covid_data$country_region == Dcountry,c('timeInt','dead')],
col='red')

```

